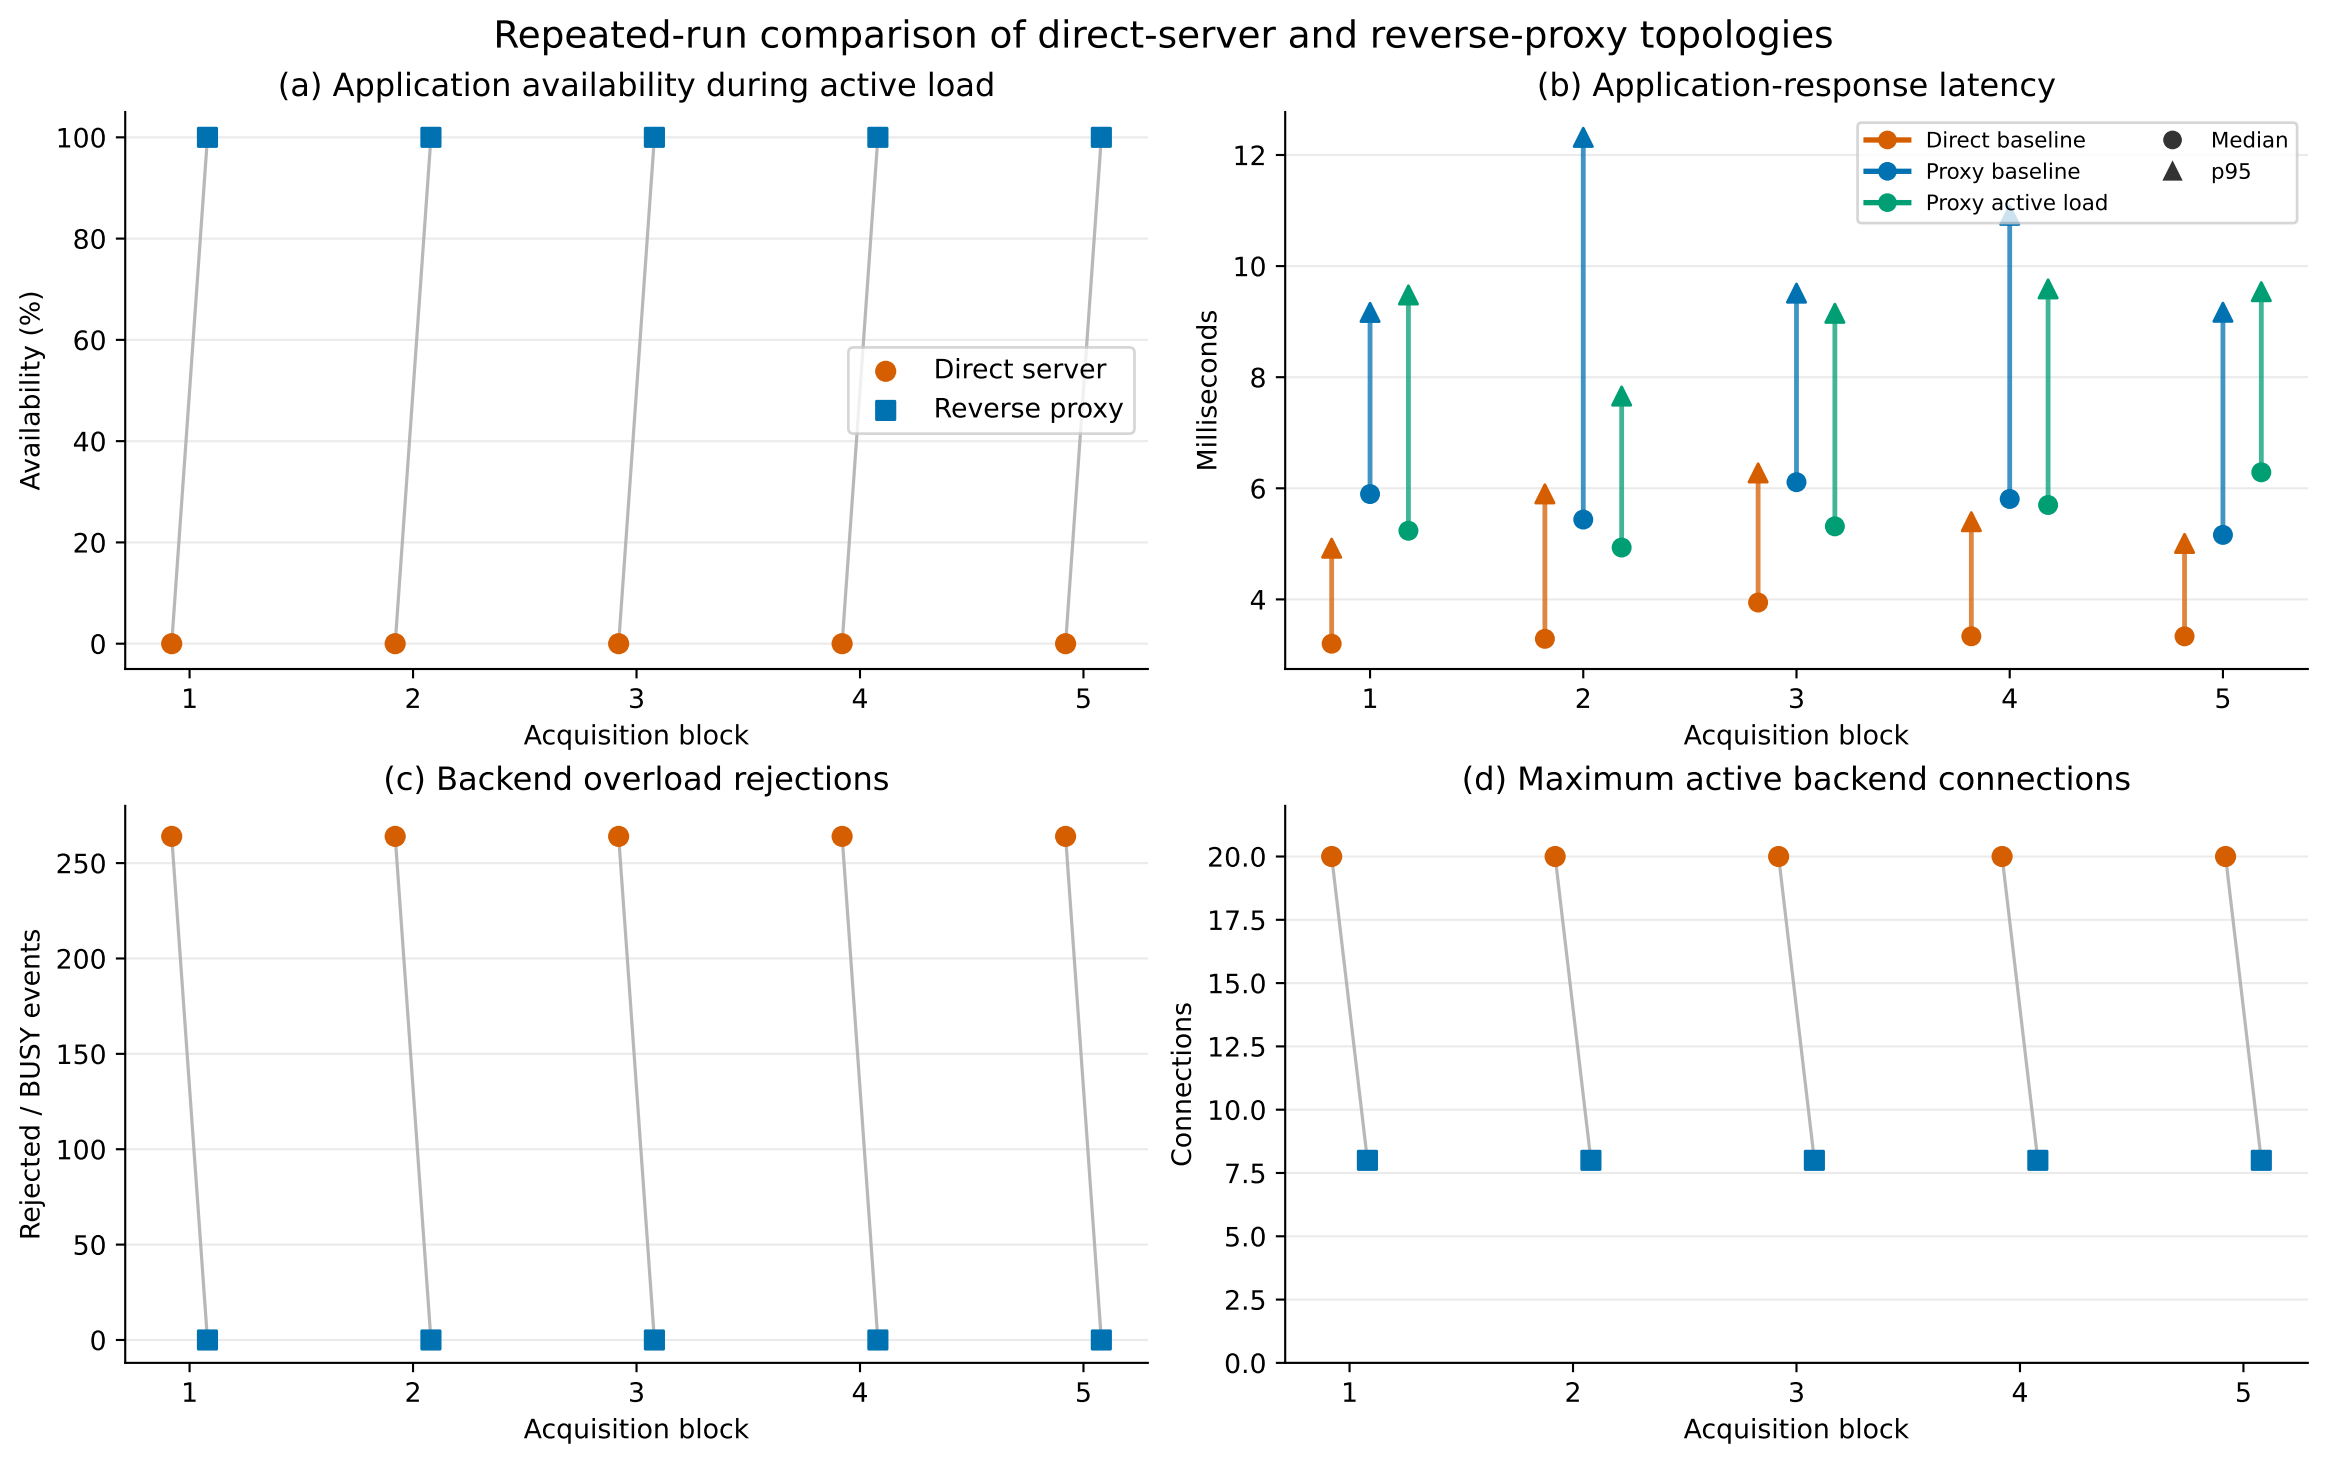
Repeated-run comparison of direct-server and reverse-proxy topologies
(a) Application availability during active load
Availability (%)
Acquisition block
Direct server
Reverse proxy
(b) Application-response latency
Milliseconds
Acquisition block
Direct baseline
Proxy baseline
Proxy active load
Median
p95
(c) Backend overload rejections
Rejected / BUSY events
Acquisition block
(d) Maximum active backend connections
Connections
Acquisition block